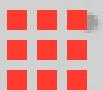
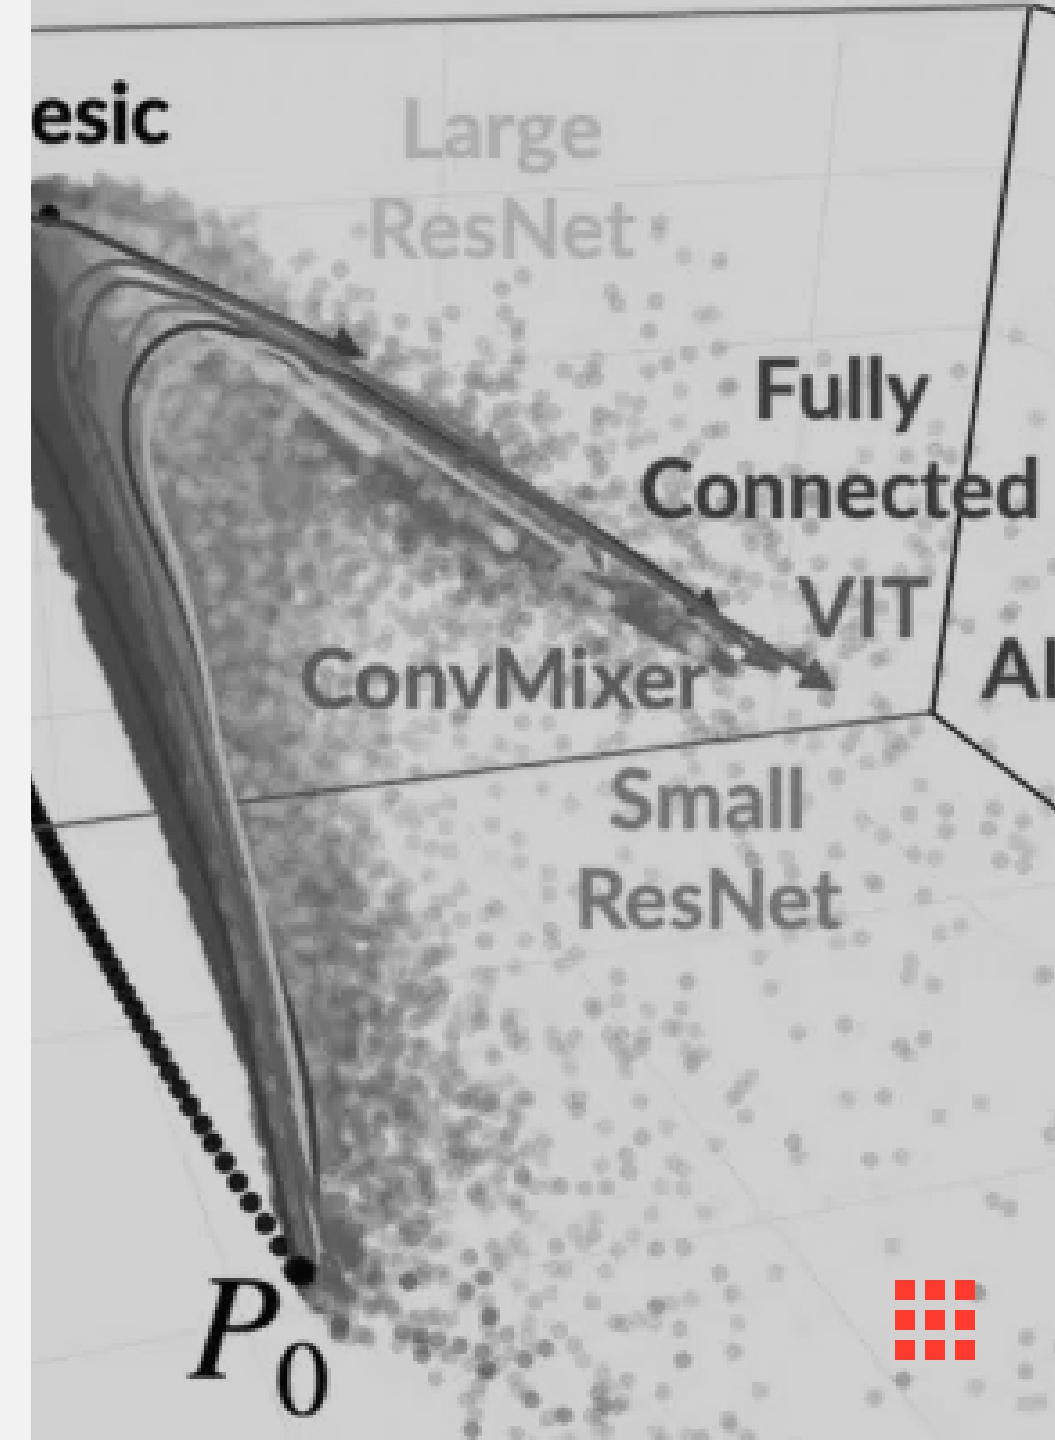


# Trajectory Geometry Determines Algorithmic Structure

A Paradigm Shift in Neural Network  
Training

BY GRISUNO



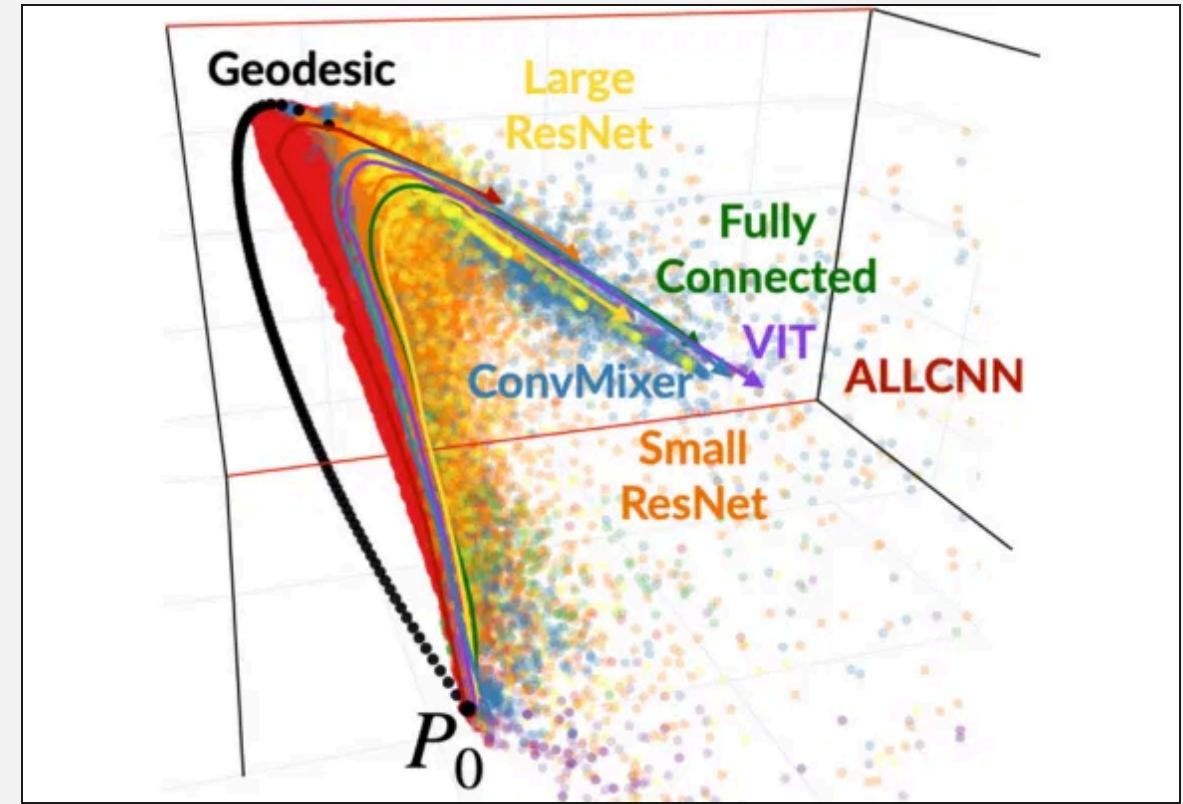
# The Central Paradigm Shift

**Traditional View:** Algorithmic structure passively emerges from optimization.

**New Paradigm:** Algorithmic structure is **actively constructed** through precise manipulation of gradient dynamics.

Controlled Trajectories  
Uncontrolled (Local Minima)  
Critical Batch Size

68% Success  
32% Failure  
[24, 128]



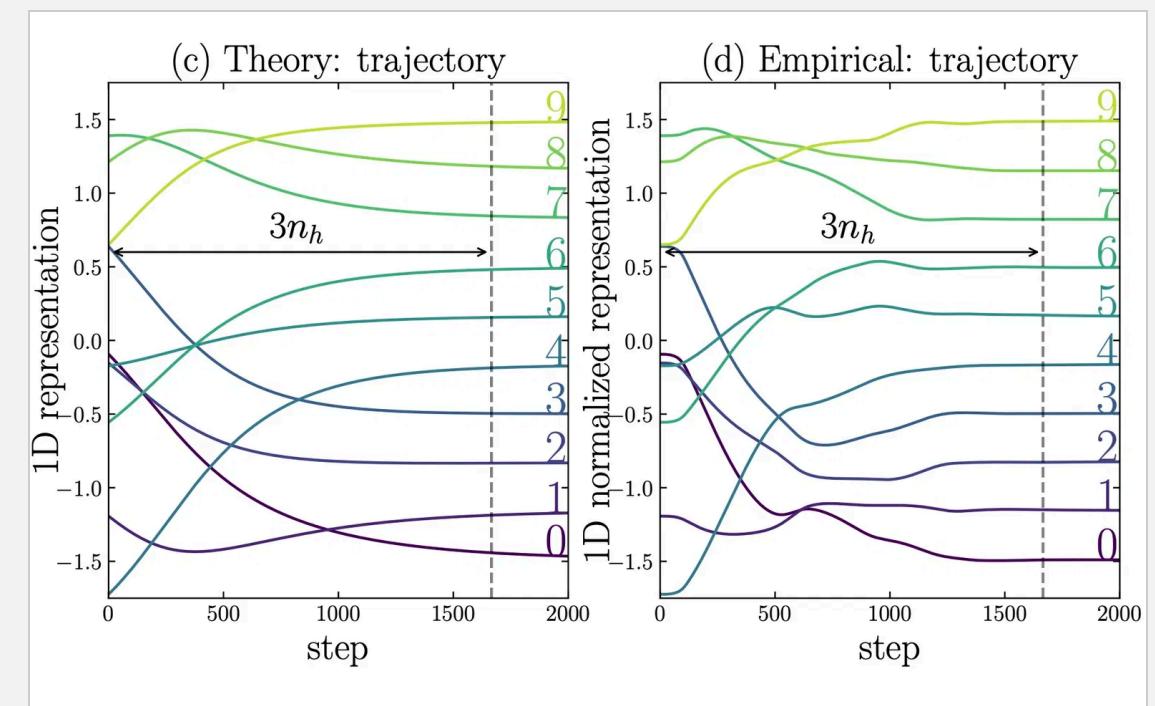
# The Grokking Problem & Verification Challenge

When a network groks, has it learned the algorithm, or has it found a local minimum that happens to generalize?

## VERIFICATION FRAMEWORK

- **Sparsification:** Prune to exact rank-7 structure (Strassen uses 7 multiplications vs. naive 8).
- **Discretization:** Project weights to integer coefficients in  $\{-1, 0, 1\}$ .
- **Zero-Shot Transfer:** Verify correctness from  $2 \times 2$  to  $64 \times 64$  matrices without retraining.

Only solutions passing all three steps represent genuine algorithmic crystallization.



Evolution of representations during training. The verification framework distinguishes true structural learning from mere statistical correlation.

# The Two-Phase Protocol for Algorithmic Induction

## 1 Training (Induction)

- **Batch Size:** [24, 128] (Critical for trajectory geometry)
- **Optimizer:** AdamW with weight decay  $\geq 1e-4$
- **Duration:** 1000+ epochs (Until grokking occurs)
- **Goal:** Reach basin of attraction

## 2 Verification (Crystallization)

- **Sparsification:** Prune to exactly 7 active slots
- **Discretization:** Round weights to  $\{-1, 0, 1\}$
- **Validation:** Verify matrix multiplication correctness

### CRITICAL INSIGHT

Phase 1 alone achieved 0% success in early experiments. Explicit sparsification (Phase 2) is required to crystallize the approximate solution into the exact algorithm.

$$\begin{aligned} p1 &= a(f - h) \\ p3 &= (c + d)e \\ p5 &= (a + d)(e + h) \\ p7 &= (a - c)(e + f) \end{aligned}$$

$$\begin{aligned} p2 &= (a + b)h \\ p4 &= d(g - e) \\ p6 &= (b - d)(g + h) \end{aligned}$$

The  $A \times B$  can be calculated using above seven multiplications.  
Following are values of four sub-matrices of result  $C$

$$\left[ \begin{array}{c|c} a & b \\ \hline c & d \end{array} \right]_X \times \left[ \begin{array}{c|c} e & f \\ \hline g & h \end{array} \right]_Y = \left[ \begin{array}{c|c} p5 + p4 - p2 + p6 & p1 + p2 \\ \hline p3 + p4 & p1 + p5 - p3 - p7 \end{array} \right]_C$$

$X$ ,  $Y$  and  $C$  are square matrices of size  $N \times N$   
 $a, b, c$  and  $d$  are submatrices of  $A$ , of size  $N/2 \times N/2$   
 $e, f, g$  and  $h$  are submatrices of  $B$ , of size  $N/2 \times N/2$   
 $p1, p2, p3, p4, p5, p6$  and  $p7$  are submatrices of size  $N/2 \times N/2$

Feature	Type
Rank-7 Constraint	Engineered
Integer Coeffs	Engineered
Zero-shot Transfer	Emergent
Discrete Values	Mixed

# The Batch Size Investigation

## 01 OBSERVATION

Batch sizes in [24, 128] achieve 68% crystallization success. Other values largely fail.

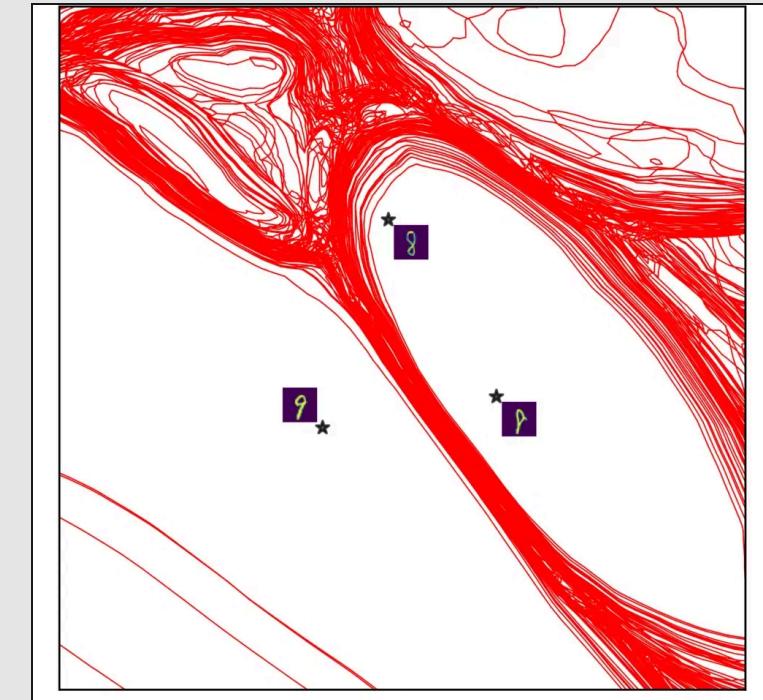
## 02 HYPOTHESIS (REJECTED)

Hardware cache effects?

**Evidence Against:** Memory analysis shows even  $B=1024$  fits comfortably in L3 cache (321 KB vs  $\geq 1\text{MB}$ ).

## 03 TRUE MECHANISM

**Gradient Covariance Geometry.** Optimal batch size stabilizes the condition number, creating trajectories where noise balances exploration and stability.



**Paradigm Demonstration:** Solutions at  $B=32$  and  $B=512$  may have identical loss, but only specific trajectory geometries reach the narrow basin containing the algorithm.

# Extreme Fragility Reveals Narrow Basins

Algorithmic solutions occupy **extremely narrow basins** in weight space.

Weights must converge precisely to integer values. Even microscopic deviations prevent structural crystallization.

## BROADER IMPLICATIONS

Reproducibility failures may reflect missed basins, not architectural flaws.

"Trajectory Engineering" is required to steer dynamics into these narrow regions.

Standard hyperparameter tuning is insufficient without geometric control.

## Noise Stability Test

NOISE ( $\Sigma$ )	SUCCESS RATE
0.001	0%
0.005	0%
0.010	0%
0.050	0%

Critical Threshold

$\sigma < 0.001$

# **Statistical Validation & Empirical Robustness**

Analysis of 195 training runs confirms that batch size is a dominant control parameter for algorithmic crystallization.

# Zero-Shot Transfer Validates Genuine Algorithmic Learning

MATRIX SIZE	MAX ERROR	STATUS
2 × 2	2.38e-07	PASS
4 × 4	1.91e-06	PASS
8 × 8	6.20e-06	PASS
16 × 16	2.15e-05	PASS
32 × 32	8.13e-05	PASS
64 × 64	2.94e-04	PASS

## WHY THIS MATTERS

Zero-shot transfer is the **definitive test**. A network that merely memorized patterns or found a convenient local minimum would fail catastrophically when expanded to larger matrices.

## CONTRAST WITH FAILED RUNS

The 32% of runs that fail verification achieve low test loss on 2×2 matrices but **cannot transfer**. They learned generalizing local minima, not the true Strassen algorithm.

# Implications for Deep Learning Reproducibility

Reframing reproducibility as a trajectory engineering problem.

## THE CHALLENGE

Many results are hard to reproduce even with identical hyperparameters.

Standard explanations like implementation details or data variations are often **insufficient** to explain why one run succeeds and another fails.

## THE INSIGHT

**Trajectory Geometry** is the missing factor.

Random initialization or hardware noise leads to different trajectories. If the target solution is in a **narrow basin**, one trajectory may reach it while another settles into a nearby local minimum.

## THE SHIFT

We must move from "Hyperparameter Tuning" to **Trajectory Engineering**.

It is necessary to understand which parameters control geometry and how to **steer trajectories** toward target basins.

# Conclusion

Active Construction Over Passive Emergence

## CORE PARADIGM SHIFT

### TRADITIONAL VIEW

Train networks, analyze final solutions, report emergent properties.



### NEW VIEW

**Engineer** training trajectories, **verify** structural properties, understand **why** specific conditions succeed.

## KEY CONTRIBUTIONS

- 02 **Verification Framework:** Explicit criteria distinguish genuine algorithmic learning from generalizing local minima.
- 04 **Empirical Evidence:** 68% success rate under controlled conditions (195 runs).
- 06 **Mechanistic Insight:** Batch size effects reflect gradient covariance geometry.
- 08 **Fragility Documentation:** 0% success with noise  $\sigma \geq 0.001$  reveals narrow basins.
- 10 **Reproducibility:** A new framework based on trajectory engineering.

Trajectory geometry determines whether networks reach algorithmic solutions.  
The 32% failure rate reflects the fundamental narrowness of algorithmic basins.

# References & Reproducibility

## KEY REFERENCES

### Deep Networks Always Grok and Here is Why

Humayun et al. (2024) • arXiv:2402.15555

### Superposition as Lossy Compression

Bereska et al. (2024) • arXiv 2024

### Algorithmic Induction via Structural Weight Transfer

grisuno (2025) • Zenodo DOI: 10.5281/zenodo.18263654

## REPRODUCTION

[github.com/grisuno/strass\\_strassen](https://github.com/grisuno/strass_strassen)

```
git clone https://github.com/grisuno/strass_strassen
cd strass_strassen
pip install -r requirements.txt
python app.py
```

## SUCCESS CRITERIA (PRE-TRAINED CHECKPOINTS)

- ✓  $\delta = 0$  (Weights exactly integers in  $\{-1, 0, 1\}$ )
- ✓ Max Error  $< 1e-5$  (Correct multiplication)
- ✓  $S(\theta) = 1$  (Successful crystallization)